

# Scaling Rails

a system-wide approach to performance optimization

MSc thesis presentation

July 2010  
MIEIC / FEUP

# Structure

# Context

## Problem

## Approach

## Results

## Conclusions

Context

Problem

Approach

Results

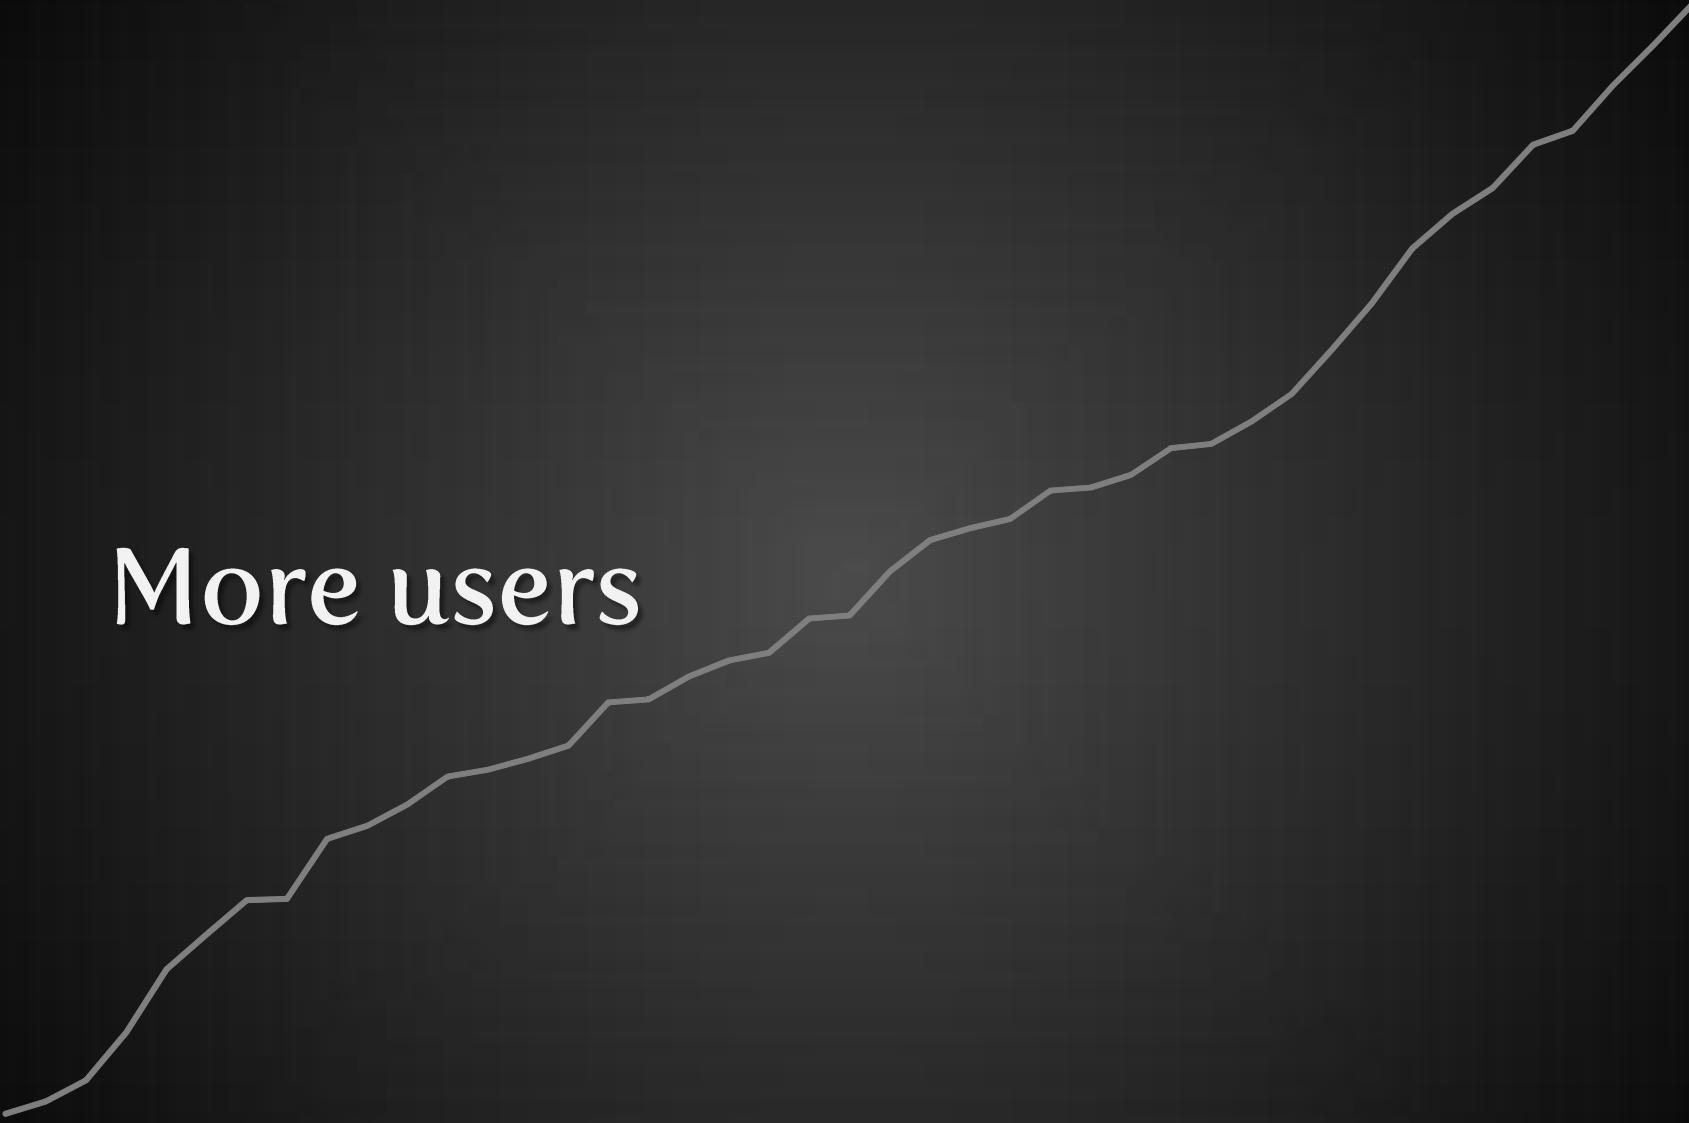
Conclusions

# Context

Problem  
Approach  
Results  
Conclusions

# Web

# More importance



More users

# Evolution

# Web 2.0

# Better UX

# Better response times

# Better web, better tools

# Ruby on Rails

# Fast development

# Full-featured applications

# Popularity

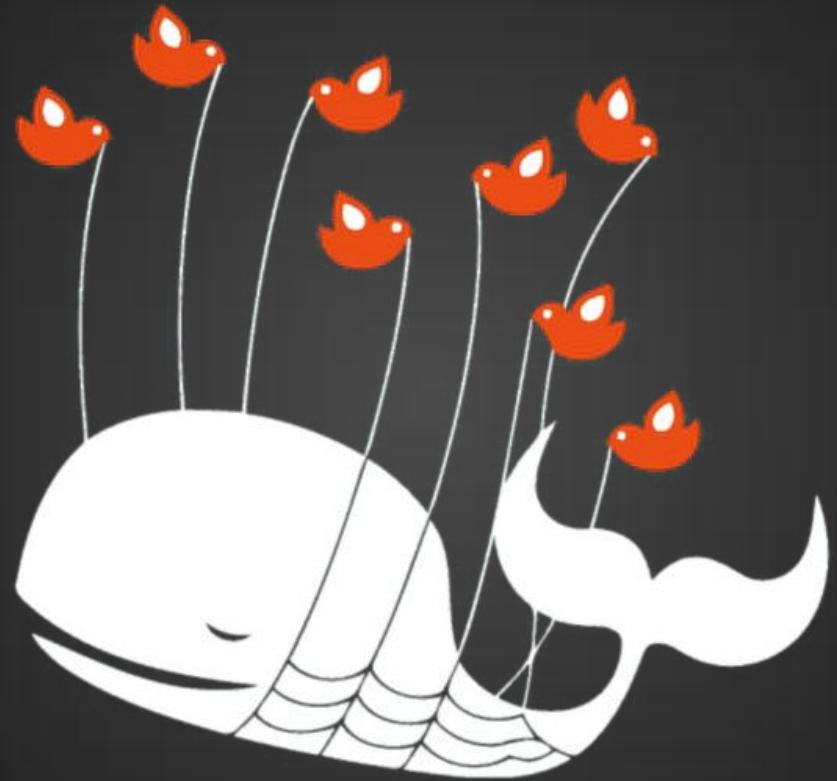
# Scalability

# Performance

# Popular websites

# Rails doesn't scale

# Twitter



# YellowPages

# “Just throw hardware at it”

# Rails doesn't scale

# Hulu

# Huge hardware infrastructure

# Basecamp

# David Heinemeier Hansson\*

\*Rails creator

# What's missing here?

Context

# Problem

Approach  
Results  
Conclusions

# Focus on

# Performance?

# Scalability?

# Complexity

# Rails

# Ruby

# Web server

# Database

# Operating System

Not much, actually

# Targeting 3 core issues

# Information

# There are benchmarks

# There are configurations

# There is information

# Scattered

# Outdated

# Specific

# Awareness

# Ruby 1.9

# Rails 3

# 2008

**They are *not* commonly used**

# Applications

# Plugins

# Gems

# Profiling

# Native tools

# Rails

# Ruby

# Which are **broken**

# Don't integrate with each other

# Text-based

Context  
Problem

# Approach

Results  
Conclusions

# Information

# General guide

# Centralized information

# Publicly accessible

# Awareness

# General guide

# Advantages of updating

# Show off

# Profiling

# Native tools

# Integrated

# More powerful

# More intuitive

# How?

# Target all system components

# Unified perspective

# Rails!

Context  
Problem  
Approach

# Results

Conclusions

# Operating System

# Development

# General benchmarking tool

# CPU-intensive

# I/O

# Workload

# Benchmarking

# General benchmark

# Ubuntu, Debian, CentOS, Gentoo

Ubuntu – CPU usage

Debian – Workload

Gentoo – CPU + I/O

CentOS – inconsistent

Ubuntu – CPU usage

Debian – Workload

Gentoo – CPU + I/O

~~CentOS~~ – inconsistent

# Not really conclusive

# Web server benchmark

$> 30$  seconds = failed request

# Requests / Concurrency

10000 / 1000

100000 / 1000

1000000 / 10000

# Ubuntu, Debian, Gentoo



# Apache / Nginx

# Default configurations

# Low or medium concurrency: OK

# High concurrency: FAIL

# Except...

# Gentoo

# Very stable

# High scalability

# What about BSD?

# Gentoo, FreeBSD



# Ruby benchmarking suite

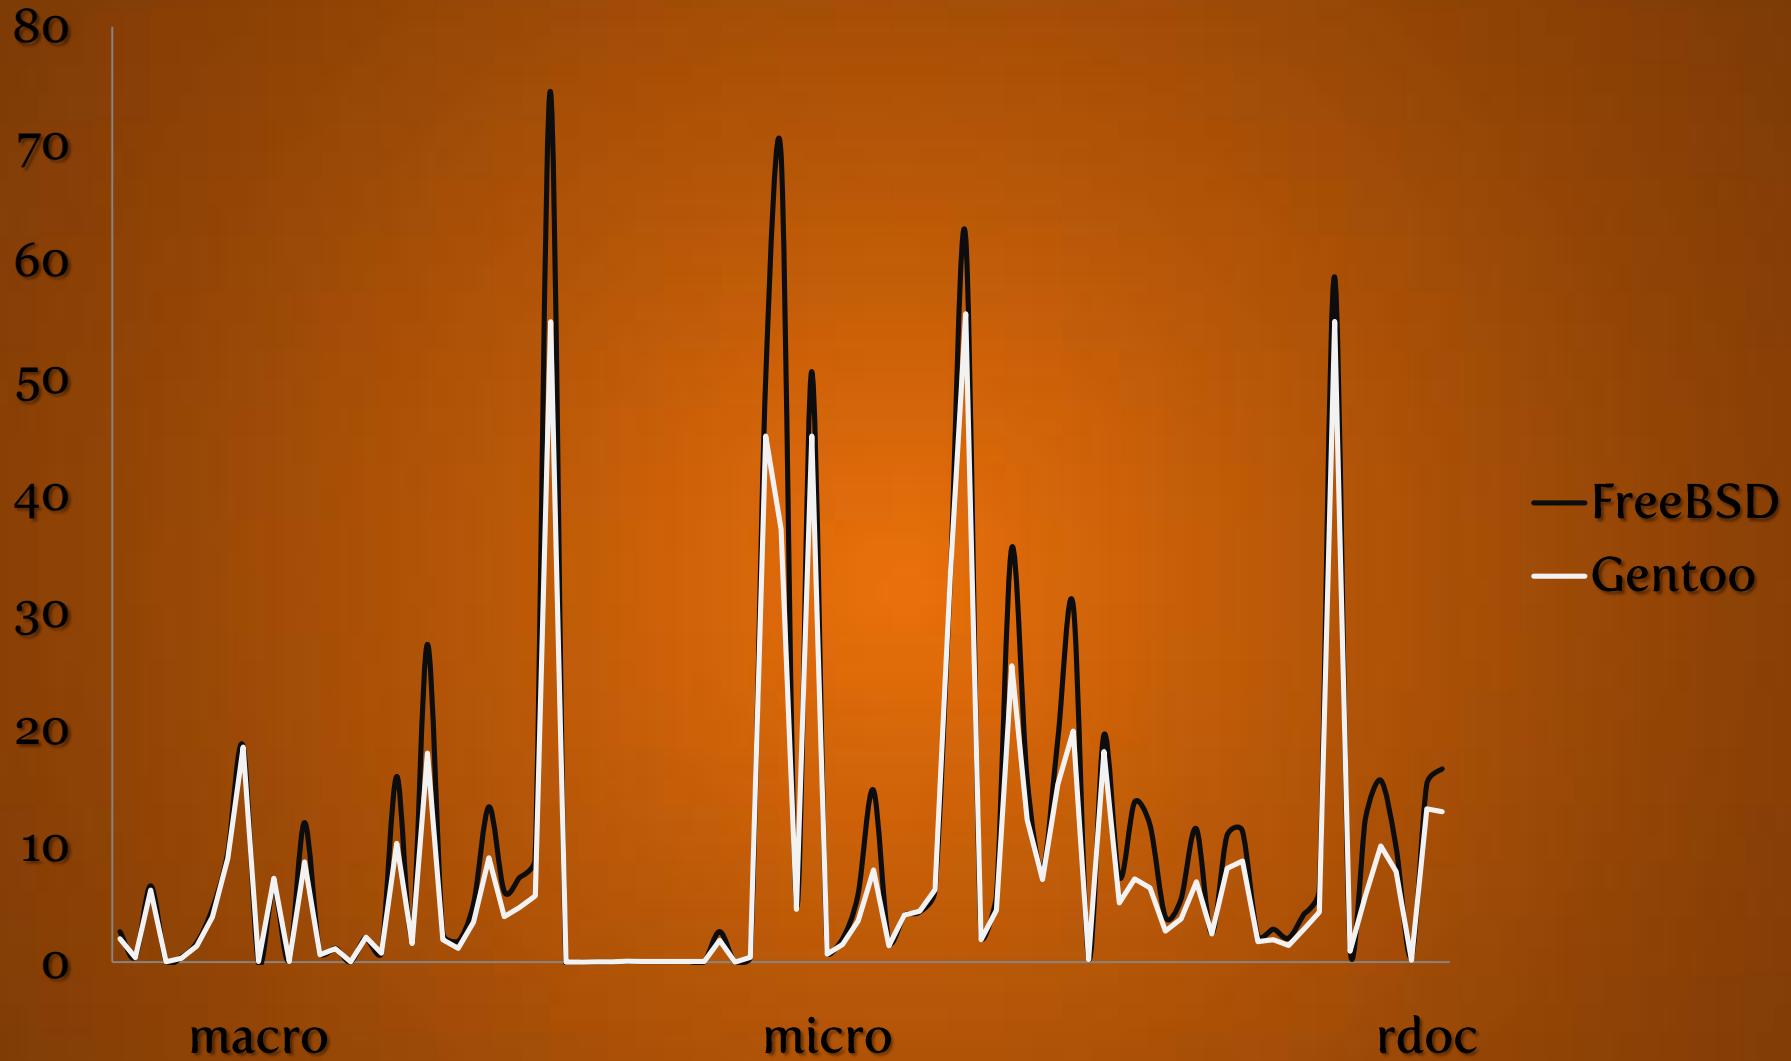
8 macro benchmarks

62 micro benchmarks

3 rdoc-related benchmarks

# Time – less is better

# MRI (1.8.7)



# More ~29% time on FreeBSD overall

# YARV (1.9.1)



# More ~22% time on FreeBSD overall

# Gentoo

# Great performance

# High configurability

# Will be used in the remaining work

# Tweaking

# Deadline I/O Scheduler

Non-preemptive configuration

100Hz timer interrupt

TCP tweaks

Kernel buffer and limit tweaks

# Operating System

Development ->

Benchmarking tool

Benchmark ->

General benchmark

Web server benchmark

Ruby benchmark

Tweaking ->

Kernel fine-tuning

# Ruby

# Benchmarking

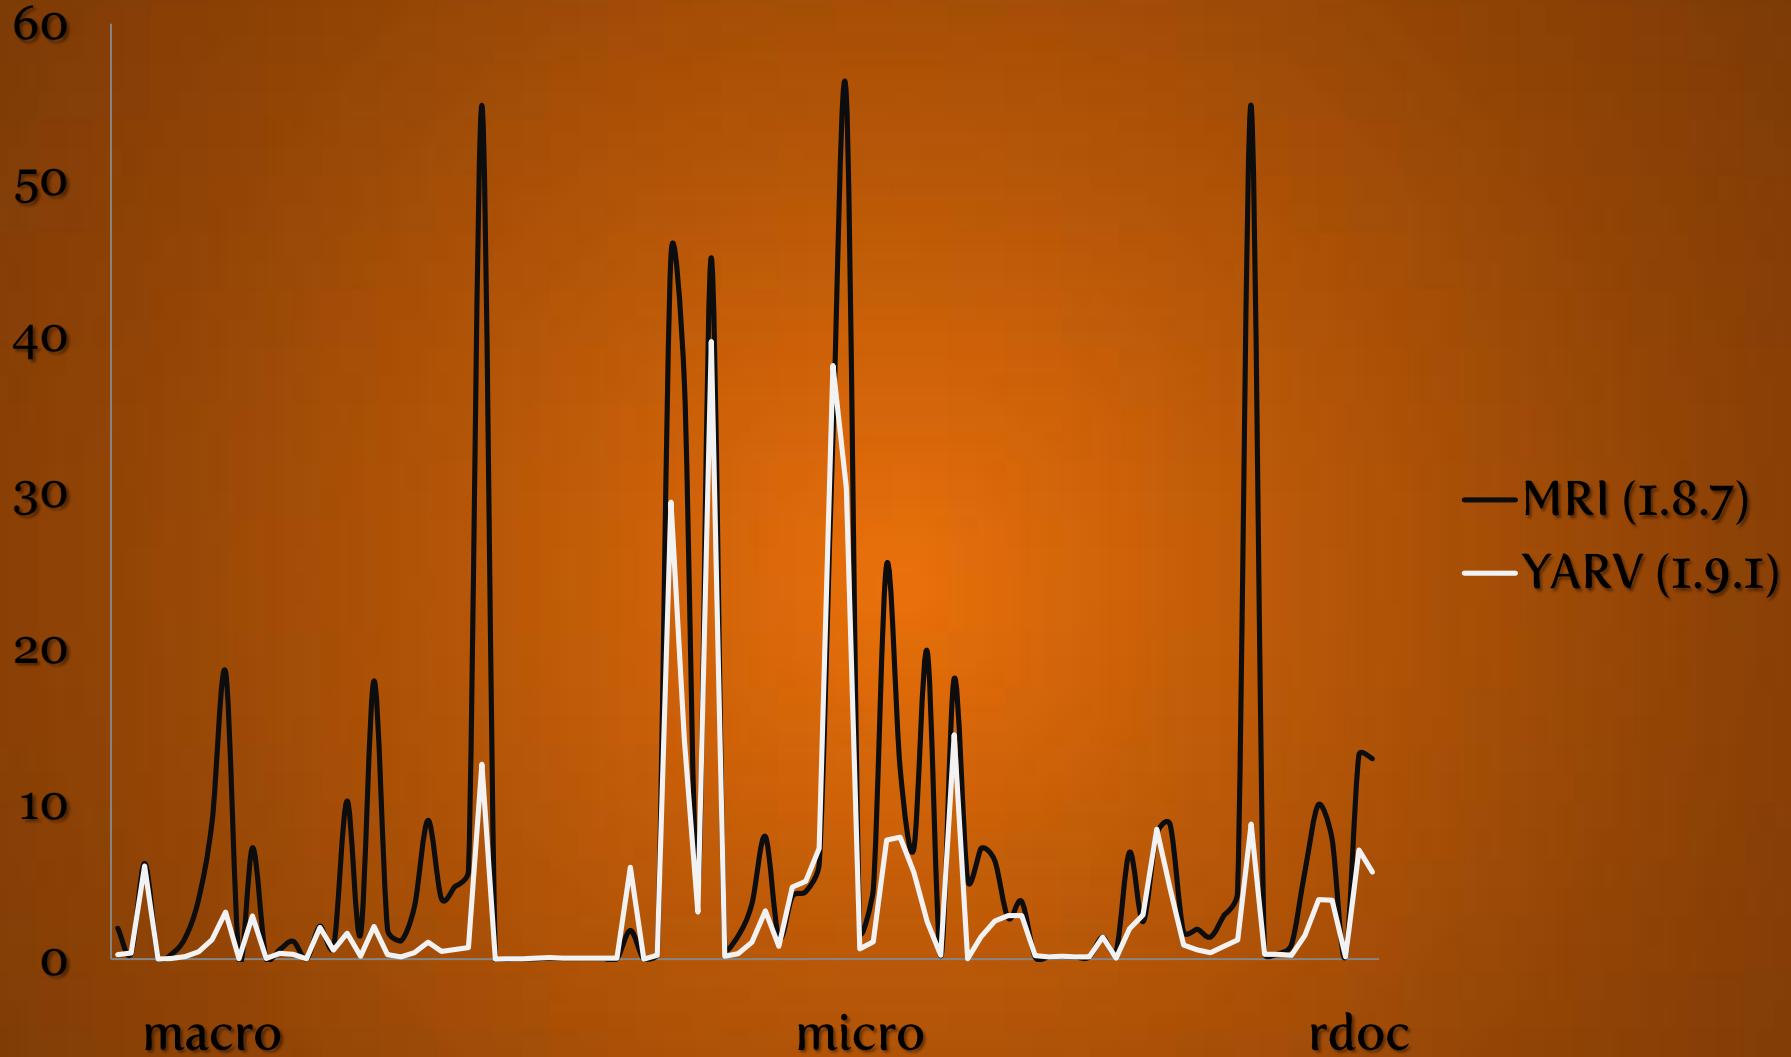
# MRI (1.8)

# YARV (1.9)

# Worth updating?

# Ruby benchmarking suite

# Ruby interpreters



# Less ~108% time on YARV overall

# YARV's performance is outstanding

# What does this mean, for Rails?

# Web server analysis will tell

# Development

# Porting *escolinhas.pt* to Ruby 1.9

*escolinhas.pt*

- > 70 models
- > 130000 lines of code
- > 40 plugins/gems
- > 10 development branches

# Rails 2.3

# Porting

# Straightforward

# Fast

# Issues

# Encoding

# Fix?

# A small patch for Rails

# Natively fixed in Rails 3

# More?

# Rare functionality changes

# Rare syntax changes

**Less than 0.1% of the time**

# Increasing YARV's flexibility

# For adaptive performance

# Tweaked MRI gave Twitter +30%

# Same principle -> YARV

# Configurable GC

# 5 configuration options

# Ruby benchmarking suite

**~15% less time needed overall**

# Slightly better memory usage

With a sample **configuration** which  
was **not fine-tunned!**

# Improving YARV's profiler

# Storing more data

# Storing correct data

# Retrieval

# In a formatted string

# Not enough

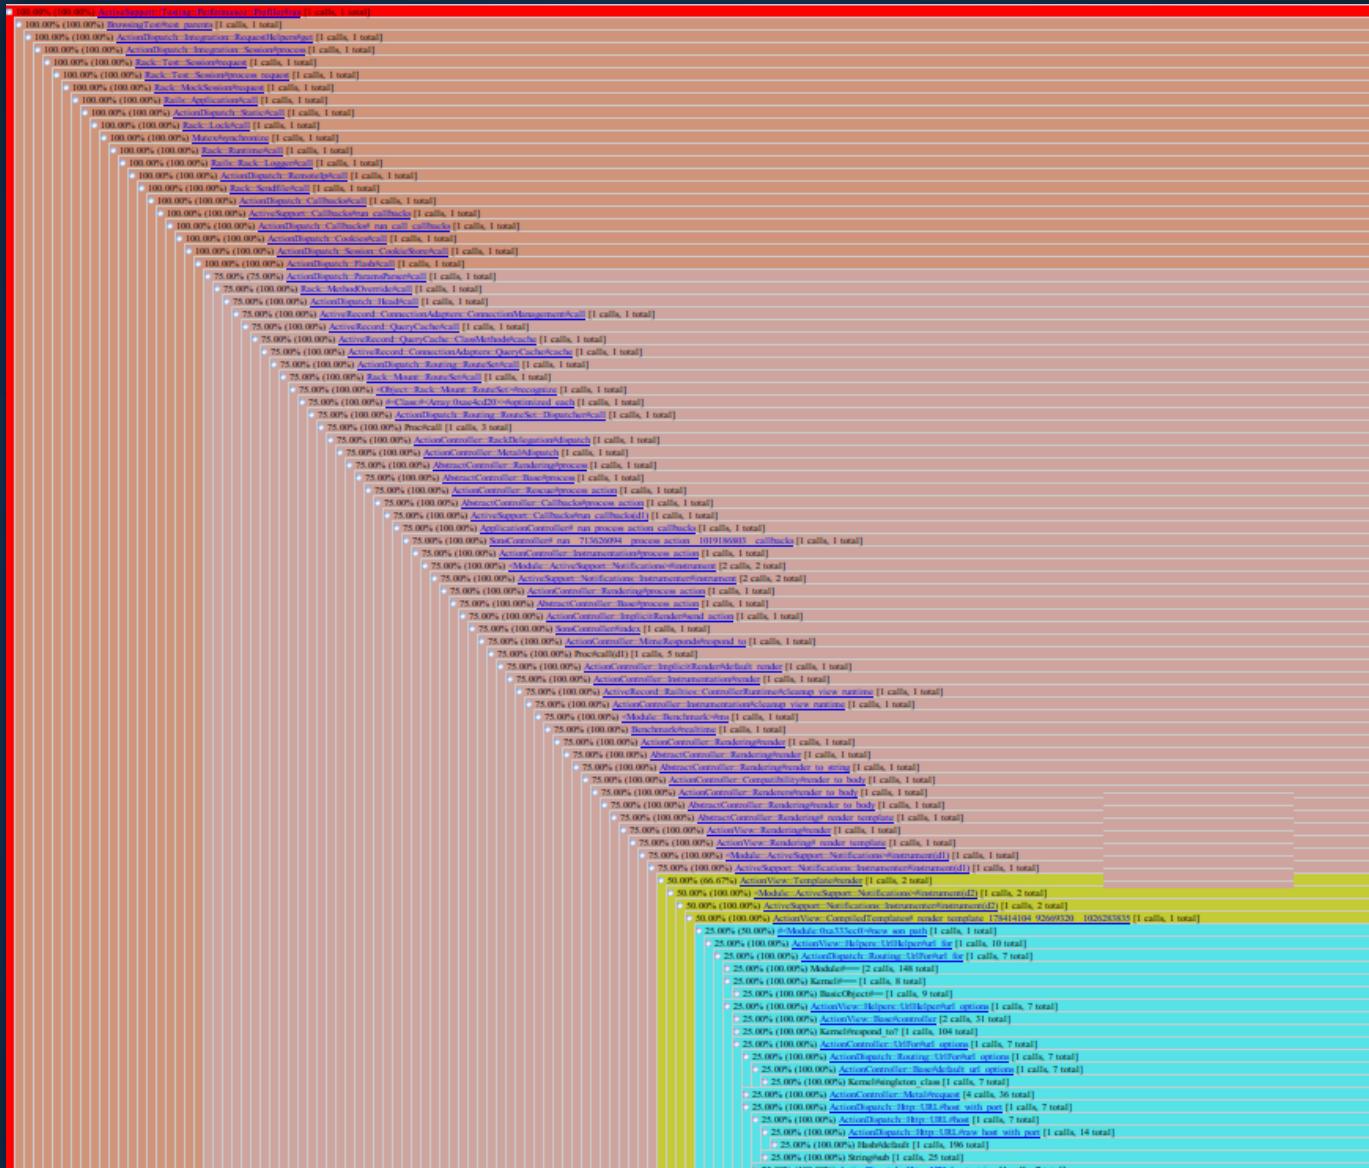
# Hash

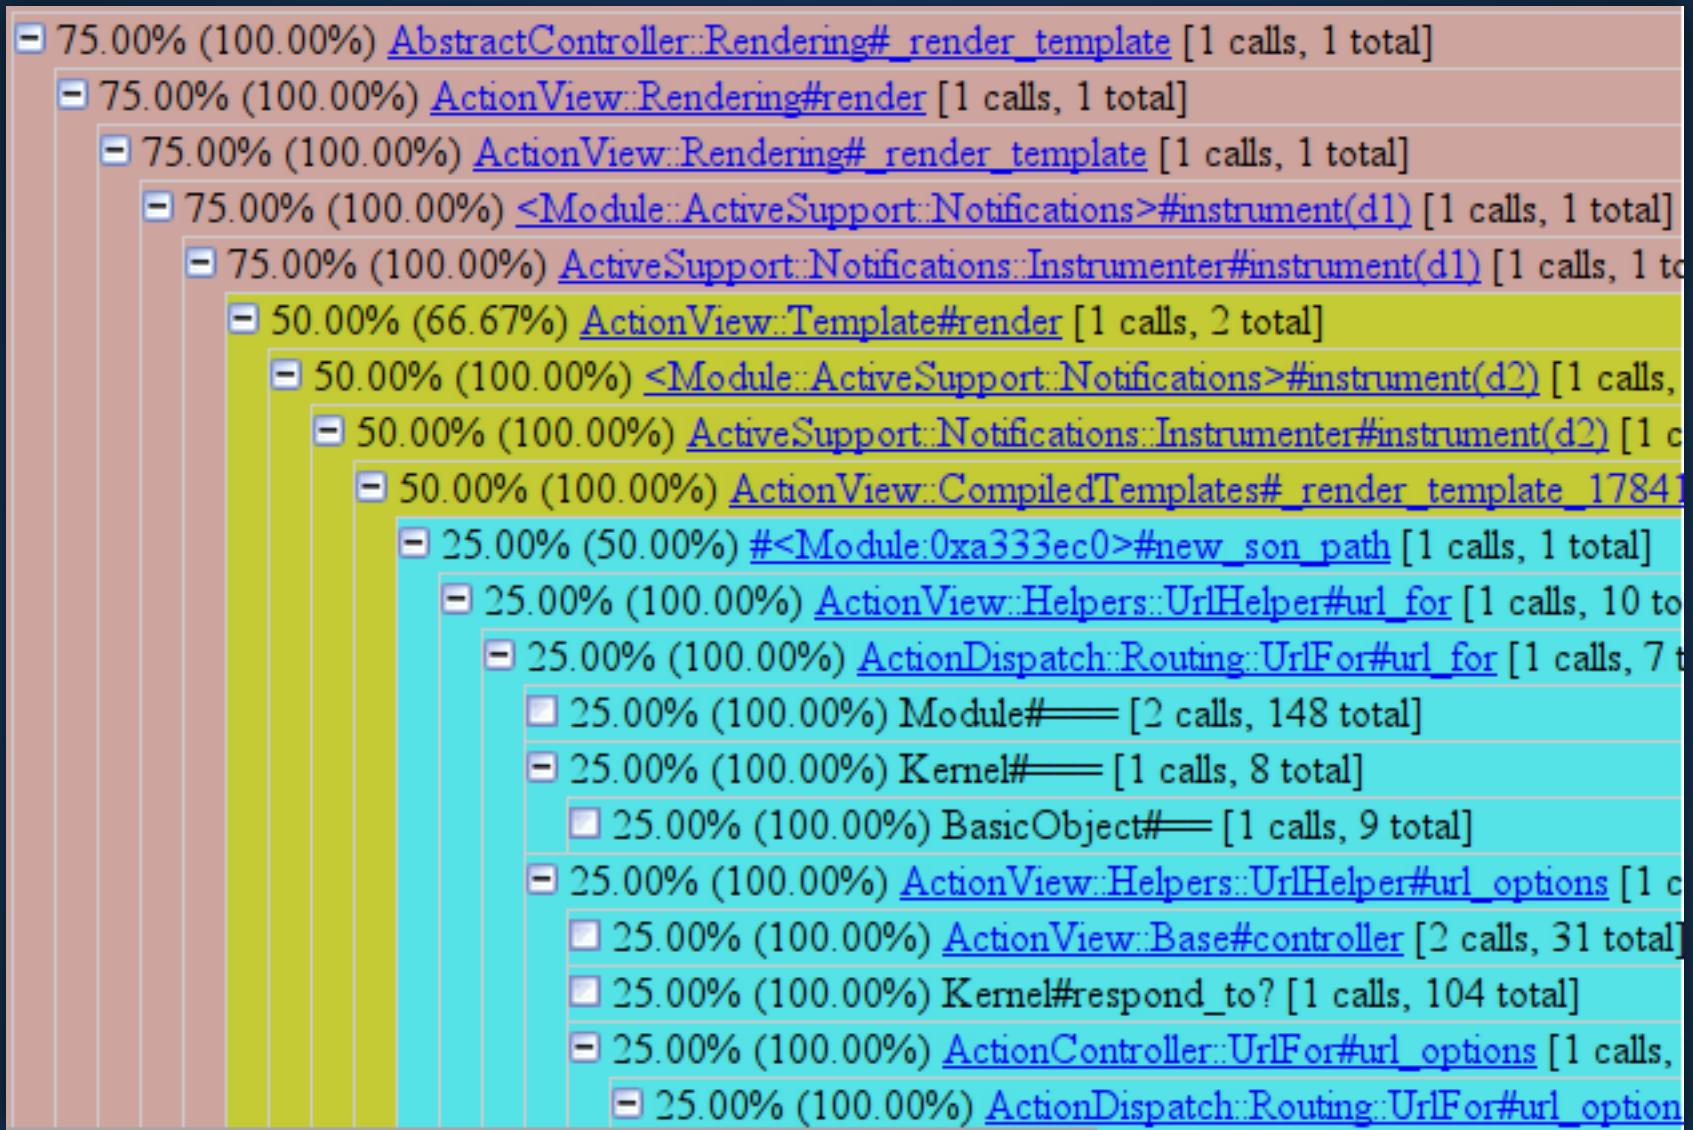
# For automated processing

# Hierarchical HTML viewer

# For human processing

And *sanity*





# Ruby

Benchmark ->

YARV and MRI

Development ->

Porting *escolinhas.pt* to Ruby 1.9

Increasing YARV's flexibility

Fixed and enhanced internal profiler

# Rails Web Servers

# Development

# Monitoring memory usage script

# Measuring consumption over time

# Very light, configurable

# Proper output format (CSV)

# Benchmarking

# Many small benchmarks

# One complete analysis



# 3 pages: heavy, medium, light

# Proxy performance

# Apache, Nginx, Cherokee

# Nginx slightly faster

# Nginx lighter

# Passenger performance

# Apache, Nginx

With Nginx it yields better results

# Thin and Unicorn comparison

# Similar performance

**Thin is slightly lighter**

# Threaded Thin performance

# Worse than non-threaded Thin

# In-depth analysis

# Passenger, Thin, Unicorn, Nginx

# 30 and 60 workers

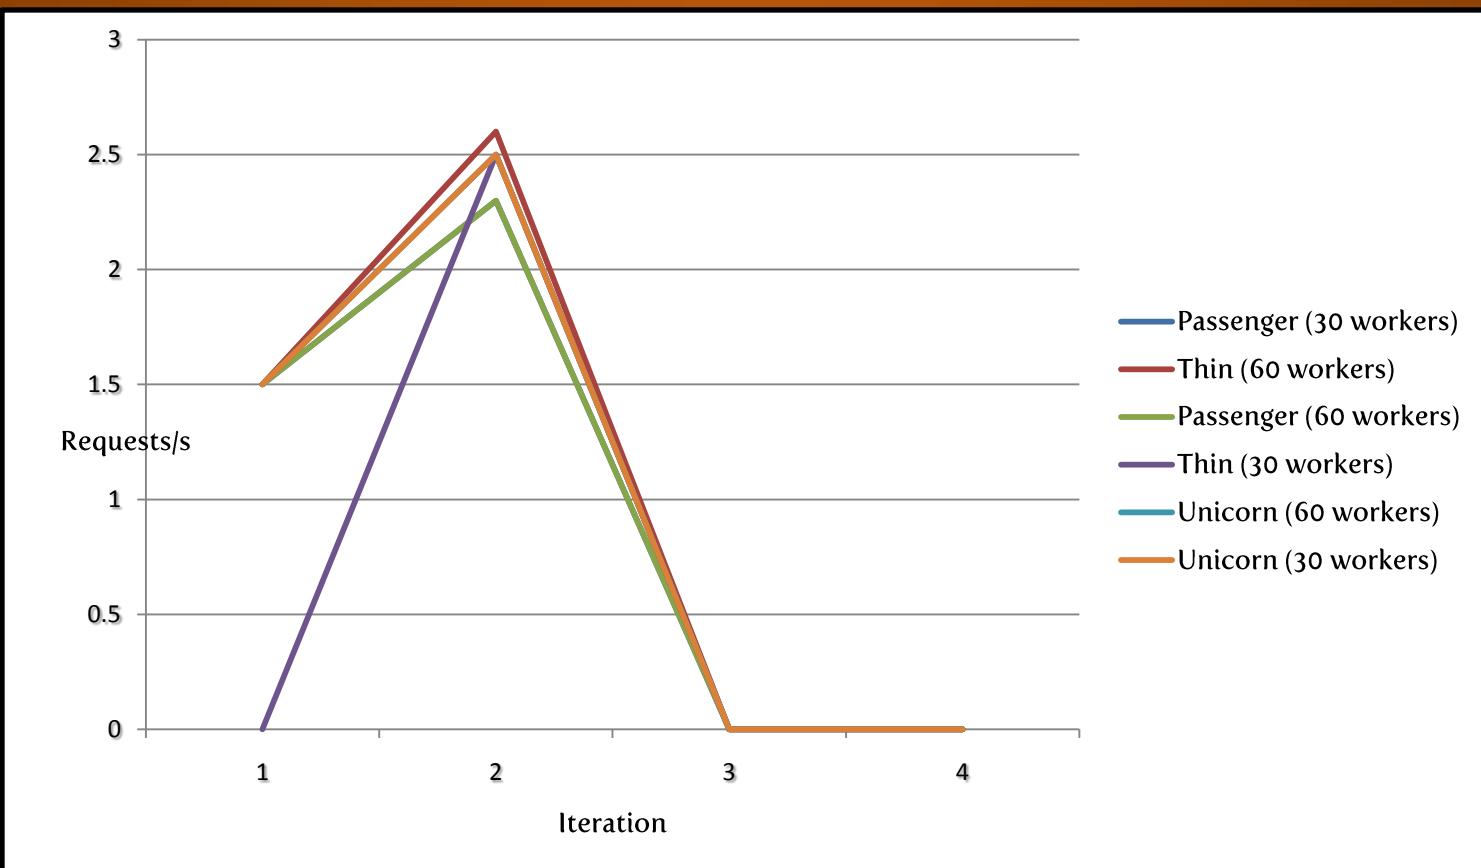
# MRI and YARV

# Using autobench

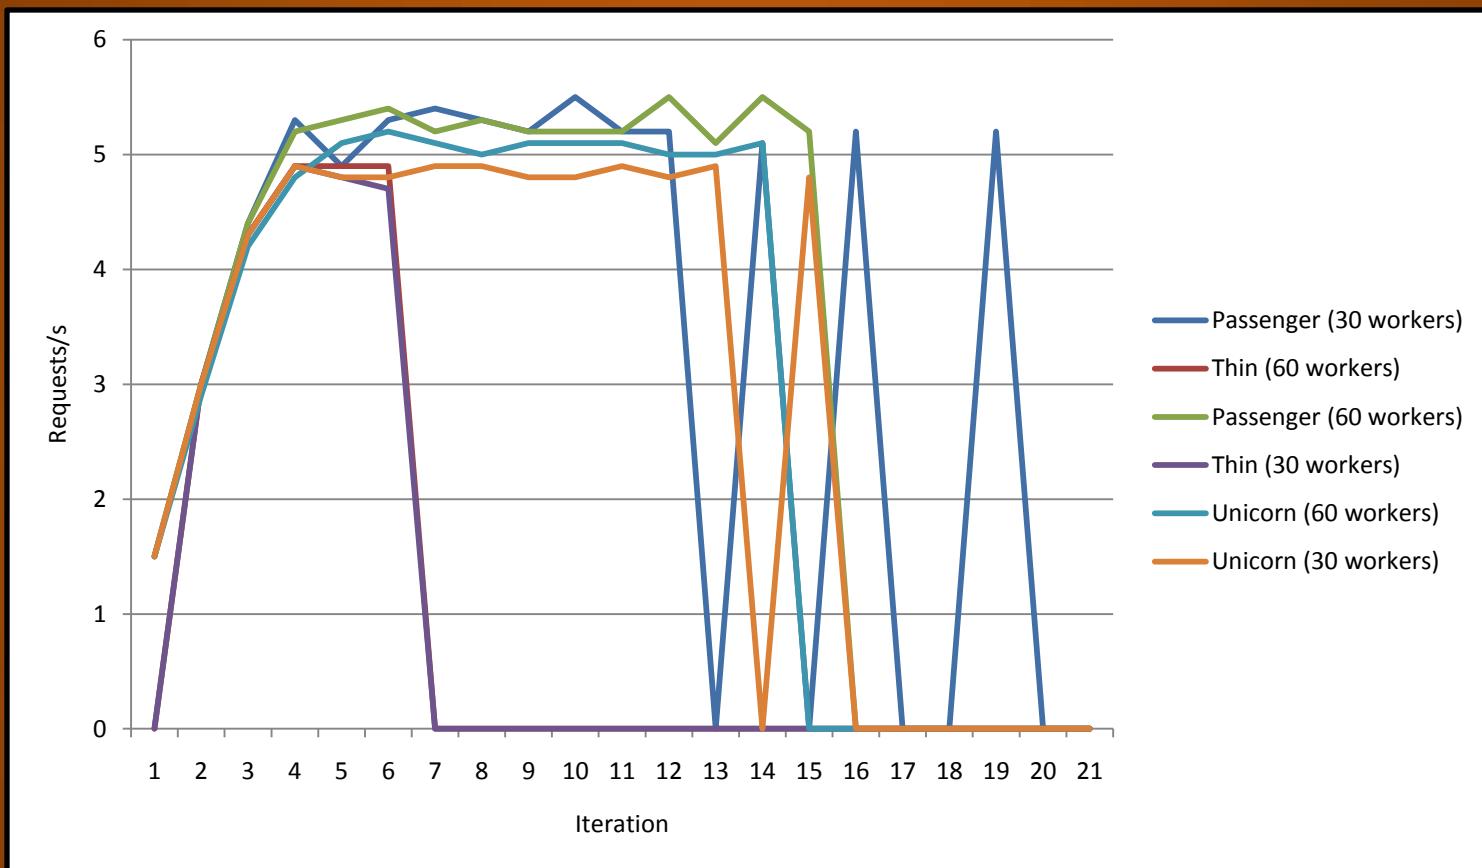
# Similar performance among web servers

# But...

# MRI (1.8)



# YARV (1.9)



1 iteration -> 15/16 iterations

2,5 requests/s -> 5 requests/s

# Huge improvement!

# Rails Web Servers

Benchmark ->

Reverse proxies

Passenger on Apache and Nginx

Thin and Unicorn

Threaded Thin

Thin, Unicorn, Passenger and Nginx

# Databases

# Development

# mysql2 adapter

# Type conversions in C

# Average 400% improvement

# Likely to replace mysql adapter

# Caveat

# More fields fetched than those used

# Solution?

# Lazy type casting

# Databases

Development ->

Adding lazy type-casting to mysql2

# Rails

# Benchmarking

# Various Rails features

Eager loading  
Transactions  
Magic finders  
Fetching in batches

# Hypothetical situations

Eager loading – 93%

Transactions – 93%

Magic finders – 68%

Fetching in batches – 10%

# Development

# Rails profiling tools

# Working

# Seamless integration with YARV

# Porting renowned Rails plugins

Used by *escolinhas.pt*

`acts_as_list`

`permalink_fu`

`acts_as_paranoïd`

# Compatible with Rails 3

# Compatible with Ruby 1.9

# Easing the porting effort

# Porting Redmine

# Compatible with Rails 3

# Compatible with Ruby 1.9

# Highly increasing the visibility of the benefits

# Ruby Summer of Code

# Benchmarking Continuous Integration

# Rails core team

# Full-stack benchmarking suite

# Automated

# Smashing performance regressions

# Community blog

# Performance-oriented article series

# Rails Magazine



# Ruby on Rails

Benchmark->

Eager loading, transactions, magic finders, fetching in  
batches

Development ->

Integrating profiling tools with YARV's  
Porting Redmine and plugins to Rails 3 and Ruby 1.9  
Benchmarking CI for Rails

Community blog

Performance-oriented article series for Rails Magazine



Context  
Problem  
Approach  
Results

# Conclusions

# General guide was created

# Information from benchmarking and analyzing

# Covering all system components

# Publicly accessible

# Profiling tools are functional

# Enhanced

# Seamlessly integrate

# Increased global awareness

# General guide

# Flexible GC in YARV

# Porting Redmine & Plugins

# Public blog

# Article series

# Official benchmarking CI

# Direct improvement

# YARV

# mysql2

# Future research

# Non-relational databases

# Native caching for Rails

# Rewriting web servers in C

# Ongoing work

# Blog

# Performance-oriented article series

# Benchmarking CI

# Thank you

Gonçalo S. Silva  
Supervisor: Ademar Aguiar (PhD.)

<http://goncalossilva.com>  
<http://twitter.com/goncalossilva>